

ClaTex: LaTeX as a Service on the Cloud

John Doe¹ and Jane Monroe²

¹Dept. Cloud, Cloud University ²Cloud National Labs

Submission Type: Research

Abstract

The time has come to offer LaTeX on the cloud.

1 Introduction

Virtualization is a basic technology for the cloud computing. The growing awareness of the advantage provided by virtualization technology is brought about by economic factors of scarce resources, government regulation and more competition.

Virtualization provides the ability of consolidation for virtual machines (VM). Multiple VMs share the same physical machine without any interference of each other. However, virtualization providers suppose VMs on the same physical machine are unlikely to use up the whole memory at the same time. So VMs are often configured more memory in total than the capacity of the physical machine. But on the worse case, VMs compete with each other leading to shortness of memory resources. Other technology like migration also provided by virtualization must start up soon. Otherwise, some VMs may suffer memory starvation and lose the performance of the application on top of that.

With the development of hardware, a typical host machine in a datacenter now packs tens to hundreds of VMs in order to maximize resource utilization. VMware ESXi hypervisor has increased the number of VMs supported per host from 32 to 1024 in recent years[10]. Memory as shared resources needs to be allocated on demand in order to full resource utilization but still the minimum QoS has to be guaranteed since users buy for it. VM providers should be able to cope with the situation when memory increases suddenly.

In virtualization environment, guest OS is transparent to the underlying hypervisor. There are only some basic monitoring statistics exposed to hypervisor such as VSZ and RSS in linux system. VSZ and RSS are coarse-grained statistics which did not reflect the memory usage in a certain period. In the past, working set theory[1] has been widely used to capture applications' memory needs. Working sets is defined as the set of all pages accessed by a process

over a given epoch. However, working set doesn't relate directly to the performance of an application. For example, if an application scans 100 pages sequentially. In a short period, the working set is 100 pages. But allocating 1 page or 100 pages gives the same memory hit ratio because the miss ratio is always 100%. MRC is the alternative to working set theory. It plots the miss rate against given physical memory. MRC may differ every second reflecting the memory requirement changes of the application which gives us a hint to dynamic allocate more or reclaim spare memory in order to utilize the host machine's resources better.

MRC theory is a general method to curve the miss rate and each part of the memory hierarchy in computer storage including cache, memory and etc. Recently many research focus on the cache MRC construction. Centaur[6] implements host-side SSD cache partition for VMs that provides both lower cache miss rate and better performance isolation and performance control for VM workloads. Like Centaur, Multi-Cache[8] presents multiple cache management on a set of storage devices and intelligently manages how caches are shared by competing VMs. Moirai[9] presents a tenant- and workload-aware system that allows data center providers to control their distributed caching infrastructure.

2 Motivation

Since the needs of oversold of memory and still high performance assurance of VMs. Memory prediction is still a valuable research area especially on the area of cloud computing. To relate the memory size and the performance, a traditional approach is to use page miss ratio curve (MRC). MRC shows the miss rate under given memory size. A way to track MRC is to use LRU stack algorithm[7]. But the time complexity and space usage is high which makes it hard to track MRC online.

Recently there are some breakthroughs to track MRC in efficient ways. Counter Stacks[13] use probabilistic counters to estimate approximate MRC at a fraction of the cost of traditional techniques. SHARDS[11] simply uses sampling to reduce the size of input trace. Both ap-

proachs reduces the time complexity and space cost in constructing MRC and makes it possible to online tracking. AET[2] describes a new kinetic model for MRC construction of LRU caches based on average eviction time(AET). AET runs in linear time asymptotically and uses sampling to minimize the space overhead.

These novel techniques track MRC in low cost. However, none of them makes it in practice on real time system especially on constructing online MRC for virtual machines in virtualization environment. MEB[12] uses optimized balancing tree to construct MRC in virtualization and dynamically balance memory allocation across all virtual machines on top of a single physical machine. However there are still 30% overhead. Intermittent Memory tracking which turns off the tracking system during steady periods cuts down the overhead to acceptable range. So online nonstop efficient MRC tracking in virtualization is still a challenge.

2.1 AET Theory

The eviction algorithm used by cache system is usually Least Recently Used (LRU). According to LRU algorithm, cache miss and cache hit both cause the cache block to move no matter how the cache is organized, LRU priority list or LRU stack[5]. AET model relates to the average eviction time of cache block. Cache block may be reused several times before it is evicted. the eviction time is the time between the last access and eviction.

AET model is build on the probability of the cache block movement. $AET(c)$ is the Average Eviction Time for all data evictions in a fully associative LRU cache of size c . Suppose Tm is the average time a cache block moves to position m . Apparently $T0 = 0$ and $AET(c) = Tc$. Suppose $rt(t)$ is the reuse numbers of reuse time t and n is the total access number. Using $f(t)$ to represent the ratio of the access number with reuse time t . $f(t) = \frac{rt(t)}{n}$. Using $P(t)$ to represent the probability of reuse time that greater than or equal to t . $P(t) = \sum_{x \geq t} f(x)$. The movement of cache block is related to the probability $P(t)$. Suppose a cache block is in position m . It will move to the next position $m + 1$ if the next access data's reuse time is greater than Tm whose probability is $P(Tm)$. Using another word, the speed of cache block in position $m(v(Tm))$ is exactly equals to $P(Tm)$. Since the integration of the speed will be the distance, we will get the following equation.

$$\int_0^{AET(c)} v(t)dt = \int_0^{AET(c)} P(t)dt = c \quad (1)$$

Given a cache of size c , according to 1, we can conduct MRC if we know the probability distribution of $P(t)$. $MRC(c) = P(AET(c))$. The probability of the access whose reuse time is greater than average eviction time of

the cache is also this cache's miss ratio. In another word, MRC can be easily calculated if we know the reuse time distribution. And reuse time distribution can be estimated by reuse time histogram.

$$MRC(m) = \frac{\sum_{i=AET(m)+1}^{\infty} rtd[i]}{\sum_{i=0}^{\infty} rtd[i]} \quad (2)$$

2.2 AET Sampling

According to AET theory, Reuse time histogram is needed to calculate the MRC. In order to get a full reuse time histogram, we need a complete access sequence. However, the overhead to get full access sequence is not acceptable. Take SPEC2006 as an example, one benchmark will have hundreds of millions of memory access. On one hand, it's hard to get full memory access. On the other hand, if every memory access should be recorded and calculate reuse time. The performance of the system for sure will be degraded severely.


AET theory presents sampling techniques. Using AET random sampling with sampling rate 1×10^{-6} , the Mean Absolute Error (MAE) is 0.01. That makes online tracking MRC possible.

3 Design

3.1 Implementation in Virtualization

In virtualization, host machine knows nearly nothing about virtual machines. Only privileged instructions are emulated to overcome the limitations arising from guest operating system running in Ring 1 and VMM running in Ring 0[3]. To illustrate how AET theory is put into virtualization, we have to understand how memory is managed in full virtualization environment.

Page tables are the data structrue used by a virtual memory system to store the mapping between virtual addresses and physical address. Traditionally, OS fully controls all physical memory space and provides a continuous addressing space to each process. But in virtualization, mutiple guest OS share the same physical memory. There should be a method to map guest OS virtual address to real machine physical memory. One way is to use shadow page table. Guest OS will maintain its own virtual memory page table in the guest physical memory frames. For each guest physical memory frame, VMM should map it to host physical memory frame. Shadow page table maintains the mapping from guest virtual address to host physical address. Page table protection helps to keep shadow page table and guest OS page table synchronization. VMM will apply write protection to all the physical frames of guest page tables, which lead the guest page table write exception and trap to VMM.



In order to track guest page access, we have to mark the page table to introduce some artificial page fault. According to The Intel 64 and IA-32 Architectures Software Developer’s Manual[4], there is reserved bit in PTE. For example, in our experiment environment, which is on intel x86-64 processors using IA-32e paging mode, there are four bits that is reserved (bit 51:48). 3.1 Reserved bits must be 0, otherwise, the entry is used neither to reference another paging-structure entry nor to map a page. So if we have PTE collections, it’s easy to track all page reference by marking reserved bits. The PTE number is a little bit huge. For a 4GB memory computer as an example, there are up to 1048576 PTEs if 4KB page used not considering huge page(2MB). An optimization is to collect page directory entry that references a page table. In x86-64 architecture, one page directory entry points to a page table that has 512 PTE. Using page directory entry(PDE) collection cuts down the space overhead to 2048 which is acceptable.

With PDE collection, we have chances to track all page reference by marking reserved bit. Once the marked page is used by guest OS, it will trap into VMM. Reserved bit should be clear and then give back the execution right to guest OS. Before giving back the execution right, VMM can record the referenced page and do lru or aet calculation. After memory access is over, mark reserved bit again. However, if all memory access is tracked, machines can hardly run. There are millions of memory ac-

Zhao[14] first presents hot set theory. They logically divide host pages into two sets, a hot page set and a cold page set. Only accesses to cold pages are trapped. Initially, all pages are marked as cold meaning all PTE's reserved bit is set. Cold pages enter hot page set after first use and clear reserved bit. Hot page set is organized as a FIFO queue. Once a page is marked as hot, its page frame number is appended to the tail of the queue. When the queue is full, the page referred to by the head entry is degenerated to the cold set, setting reserved bit again.

3.3 Sampling

These two works of sampling have some distinctions. Shards concept is pretty simple - for each referenced location L, the decision of whether or not to sample L is based on that $\text{hash}(L)$ satisfies some condition - $\text{hash}(L) \bmod 100 < K$ means $K\%$ sampling rate. AET sampling

according to probability choose whether or not to sample a reference. It means the distance between two adjacent monitoring points is a random value. Unlike shards' method, every reference in the trace has a chance to be sampled. These two works both show good results.

In our situation, random sampling seems impossible. There is so huge overhead to get full reference to memory, not saying randomly pick some to track. What we have is a collection of PTEs. A practical way to use sample method to track page reference is randomly choosing PTE according the hash value of PTE address which is like Shards. To avoid tracking fixed set of PTE, we periodically change the sampling collection. Every period, we scan the PTE collection and randomly pick PTE to set reserved bit. In this way, we guarantee our tracking set is random and often changing.

3.4 Dynamic sampling rate

Hot page set and sampling reduce the overhead drastically. But there is still a problem. Hot page set size is based on locality. Programs have a different degree of locality. Even within one program, locality varies frequently. Fixed hot page set size and sampling rate may not reduce the overhead to the acceptable range. Meantime, if a program's memory needs are not so big, with a large hot page set size and a small track rate, the number of tracked page becomes too small. AET algorithm is based on probability. Too small samples will also affect accuracy. xxxxx reflects the overhead with fixed hot page set size and sampling rate.

xxxxxxx L

Increasing hot page set size or decreasing sampling rate both reduce overhead. How to make a choice, there is no theory support nor prior knowledge. It's a two variables problem. By fixing one variable, we change another one to watch the overhead change. Experiments result show that decreasing sampling rate reduce the overhead immediately, and increasing hot page set may or may not reduce the overhead. For example, xxxxxxxx. We run this benchmark again to check why increasing hot page set doesn't reducing overhead immediately and count the page fault number that we manual make.

xxxxxxx page fault

In the tables xxxxxxx, it shows increasing hot page set size may not decreasing page fault number. The reason is that if a program's spatial locality is much larger than our hot page set size or much smaller than that, the effect of increasing hot page set size or decreasing is not evident. Keep increasing hot page set size obvious can reduce page fault numbers but this process is slow. However changing sampling rate directly reduces the size of tracking page set. Obviously it's effect is remarkable. Since we can not slow down the performance of the vir-

tual machine, we need a mechanism to reduce overhead quickly. Our research is focused on changing sampling rate. Before that, it's important to verify whether changing hot page set size or sampling rate influence accuracy. From our experiment, these changes have little effect on accuracy. And from prior work, dynamic changing hot page set size is a method to reduce large working set program overhead[12], One ten thousandth sampling rate still maintains high accuracy[2]. Based on the prior research and our experimental verify, our work use dynamic hot page set size and sampling rate to reduce overhead.

xxxxxx graph

4 Conclusion

References

- [1] P. J. Denning. The working set model for program behavior. *Communications of the Acm*, 26(1):43–48, 1968.
- [2] X. Hu, X. Wang, L. Zhou, Y. Luo, C. Ding, and Z. Wang. Kinetic modeling of data eviction in cache. In *2016 USENIX Annual Technical Conference (USENIX ATC 16)*, pages 351–364, Denver, CO, 2016. USENIX Association.
- [3] P. M. Humble Devassy Chirammal and A. Vettathu. *Mastering KVM Virtualization*. Packt, 2016.
- [4] Intel. *Intel 64 and IA-32 Architectures Software Developer's Manual*. Intel Corporation, 2016.
- [5] C. Jo, E. Gustafsson, J. Son, and B. Egger. Efficient live migration of virtual machines using shared storage. In *ACM Sigplan/sigops International Conference on Virtual Execution Environments*, pages 41–50, 2013.
- [6] R. Koller, A. J. Mashtizadeh, and R. Rangaswami. Centaur: Host-side ssd caching for storage performance control. In *IEEE International Conference on Autonomic Computing*, pages 51–60, 2015.
- [7] R. L. Mattson, J. Gecsei, D. R. Slutz, and I. L. Traiger. Evaluation techniques for storage hierarchies. *Ibm Systems Journal*, 9(2):78–117, 1970.
- [8] S. Rajasekaran, S. Duan, W. Zhang, and T. Wood. Multi-cache: Dynamic, efficient partitioning for multi-tier caches in consolidated vm environments. In *IEEE International Conference on Cloud Engineering*, pages 182–191, 2016.

- [9] I. Stefanovici, E. Thereska, G. O'Shea, B. Schroeder, H. Ballani, T. Karagiannis, A. Rowstron, and T. Talpey. Software-defined caching: managing caches in multi-tenant data centers. In *ACM Symposium on Cloud Computing*, pages 174–181, 2015.
- [10] vSphere 6.0. Configuration maximums. <https://www.vmware.com/pdf/vsphere6/r60/vsphere-60-configuration-maximums.pdf>, 2017.
- [11] C. A. Waldspurger, N. Park, A. Garthwaite, and I. Ahmad. Efficient mrc construction with shards. In *Usenix Conference on File and Storage Technologies*, pages 95–110, 2015.
- [12] Z. Wang, X. Wang, F. Hou, Y. Luo, and Z. Wang. Dynamic memory balancing for virtualization. *Acm Transactions on Architecture & Code Optimization*, 13(1):2, 2016.
- [13] J. Wires, S. Ingram, Z. Drudi, N. J. A. Harvey, and A. Warfield. Characterizing storage workloads with counter stacks. In *Usenix Conference on Operating Systems Design and Implementation*, pages 335–349, 2014.
- [14] W. Zhao, Z. Wang, and Y. Luo. Dynamic memory balancing for virtual machines. *Acm Sigops Operating Systems Review*, 43(3):37–47, 2009.